


The AI Model Car Show


Which AI model is your Ferrari — and which is your Toyota? A performance-vs-efficiency breakdown of every major model across Anthropic, Google, and OpenAI.



Every AI model has a car equivalent.

Performance = accuracy, benchmark scores, hallucination reliability. Efficiency = token cost, speed, mileage per dollar. A Ferrari burns fuel for lap times. A Prius is built for commuting. Know which one your workload needs.


ECONOMY & EFFICIENCY



Toyota Prius

Haiku 4.5


\$1 / Mtok



Civic Type R

Gemini 3.5 Flash


\$1.50 / Mtok



Porsche 911

Sonnet 4.6


\$3 / Mtok



Ferrari F8

Opus 4.7 / GPT-5.5

\$5 / Mtok



Formula 1

GPT-5.5 Pro

\$30 / Mtok

PERFORMANCE METRIC

Hallucination rate, SWE-bench, GPQA Diamond — how accurate and reliable the output is

EFFICIENCY METRIC


\$ per million tokens in/out — how much fuel you burn per mile of useful output

THE RULE

Route 80% of traffic to Sonnet/Flash (Porsche/Civic). Reserve Ferrari for your 10 hardest problems.

Anthropic — The Reliability Team

Best hallucination rates industry-wide · MCP inventor · 4.2M weekly Claude Code developers



FERRARI F8

Claude Opus 4.7

The peak performer. When accuracy is non-negotiable and cost is secondary.

HALLUCINATION RATE

36%

SWE-BENCH PRO

64.3% #1

INPUT / OUTPUT

\$5 / \$25 Mtok

CONTEXT WINDOW

1M tokens

Fuel warning: New tokenizer inflates effective cost up to 35%. Gate behind cost guardrails — only route here for your hardest 10% of tasks.



PORSCHE 911

Claude Sonnet 4.6

The production default. High performance, daily-driveable cost. Routes 80% of workloads perfectly.

HALLUCINATION RATE

Excellent

SWE-BENCH VERIFIED

79.6%


INPUT / OUTPUT

\$3 / \$15 Mtok

SPEED VS OPUS

~2× faster

Recommended default for 80% of production workloads. Best accuracy-to-cost ratio in the Anthropic family.



TOYOTA PRIUS

Claude Haiku 4.5

Max mileage. Route all routing, classification, and extraction here — and cut your inference bill by 30–60%.

HALLUCINATION RATE

Excellent

BEST USE CASE

Routing · Classification

INPUT / OUTPUT

\$1 / \$5 Mtok

SPEED

Fastest in family


Bulk route here. Intelligent routing to Haiku for low-complexity tasks cuts inference cost 30–60% in mixed workloads.

Google — The Price-Performance Engineers

Best pure reasoning (GPQA 94.3%) · Only on-device AI · Veo 3.1 video leader · 10-country data residency

Gemini 3.1 Pro Preview

Best price-to-reasoning ratio at the frontier. Half the cost of Opus 4.7, better on pure reasoning benchmarks.



BMW M3

HALLUCINATION RATE

50% (moderate)

GPQA DIAMOND

94.3% — #1

INPUT / OUTPUT

\$2 / \$12 Mtok


CONTEXT WINDOW

1M (2M practical)

Still "Preview" status as of May 2026. Do not anchor production SLAs to preview models. GA expected "soon" per Google.

Gemini 3.5 Flash

Launched May 19, 2026. New price-performance leader for high-volume production — sporty, economical, and fast.



HONDA CIVIC TYPE R

HALLUCINATION RATE

Moderate

MCP ATLAS SCORE

83.6%

INPUT / OUTPUT

\$1.50 / \$9 Mtok

SPEED VS 3.1 PRO

~4x faster

New default recommendation for Google high-volume pipelines. Beats 3.1 Pro on coding benchmarks at 75% of the cost.




Gemini Nano 4 — The Off-Road Electric

On-device SLM via Android AICore. NPU-accelerated. \$0 inference cost. 140+ languages. The only frontier vendor with a production on-device stack — OpenAI and Anthropic are both cloud-only.

\$0
inference cost
GOOGLE ONLY

OpenAI — Raw Power, Handle With Care

Best terminal agents & computer use · Highest hallucination rate at frontier · Most expensive per Mtok



LAMBORGHINI HURACÁN

GPT-5.5

Raw power for terminal agents and computer use. Wins the benchmarks it was built for — and misfires badly on everything else.

HALLUCINATION RATE	86% — highest	TERMINAL-BENCH 2.0	82.7% — #1
INPUT / OUTPUT	\$5 / \$30 Mtok	OSWORLD	78.7% — #1

86% hallucination risk. Deploy only where grounding and citation guardrails exist. Never use for knowledge work without RAG. Price doubled vs GPT-5.4.



FORMULA 1 CAR

GPT-5.5 Pro

Extreme performance, extreme cost. Built for one specific race: frontier web research and impossible math. Useless outside those lanes.

HALLUCINATION RATE	Better than 5.5	BROWSECOMP	90.1% — #1
INPUT / OUTPUT	\$30 / \$180 Mtok	BEST FOR	Frontier research · Hard math

6× more expensive than GPT-5.5 standard. Reserve strictly for BrowseComp-level research or AIME-grade mathematics — nothing else justifies the cost.



Sunset alert: Sora 2 API and model aliases retire September 24, 2026. OpenAI Assistants API retires August 26, 2026. Migrate to Responses API now. For video, migrate to Google Veo 3.1.

SUMMARY

Who wins which race?

No vendor sweeps. The model layer has trifurcated by job. Buy the right car for the right road.

CODING ACCURACY

Anthropic wins

Claude Opus 4.7 · SWE-bench Pro 64.3% — 6 points ahead of GPT-5.5

PURE REASONING

Google wins

Gemini 3.1 Pro · GPQA Diamond 94.3% at half Opus 4.7's price

VIDEO GENERATION

Google wins

Veo 3.1 — cinematic 1080p + native audio · Sora 2 sunseting Sep 2026

ON-DEVICE / MOBILE

Google wins — uncontested

Gemini Nano 4 · \$0 inference · Only frontier vendor with on-device stack

WEB RESEARCH BREADTH

OpenAI wins

ChatGPT Deep Research · broadest 50+ source coverage

TERMINAL & COMPUTER USE

OpenAI wins

GPT-5.5 + Codex · Terminal-Bench 82.7% · OSWorld 78.7%

HALLUCINATION RELIABILITY

Anthropic wins

Opus 4.7 at 36% · Google 50% · OpenAI 86%

TOKEN COST EFFICIENCY

Google wins

Gemini 3.5 Flash at \$1.50/\$9 · Haiku 4.5 at \$1/\$5 for routing

ENTERPRISE + HIPAA

Tied: Anthropic + Google

Both offer HIPAA BAA (Bedrock / Vertex), ZDR, ISO 27001/42001

DOCUMENT RESEARCH (GROUNDED)

Google wins

NotebookLM · best citation fidelity · Audio + Video overviews

THE 3 RULES OF AI MODEL SELECTION

Rule 1: Route 80% of traffic to Sonnet 4.6 or Gemini 3.5 Flash (Porsche / Civic). They cover most jobs at a fraction of Ferrari cost.

Rule 2: Gate Opus 4.7 and GPT-5.5 Pro behind explicit cost guardrails. Only deploy for your hardest 10% of tasks.

Rule 3: Never use GPT-5.5 for knowledge work without RAG. 86% hallucination means it sounds confident and wrong — often.