

AI TOOLKIT COMPARISON · MAY 2026 · BASED ON LIVE BENCHMARKS

Anthropic vs Google vs OpenAI

Full Portfolio by Use Case

Which tools win for each job — design, coding, video, research, agents, enterprise. Includes hallucination rates, token costs, and the performance-vs-efficiency trade-off.

- Anthropic / Claude
- Google / Gemini
- OpenAI / ChatGPT

- Use Cases
- Performance vs Efficiency
- Token Cost Table
- Winner Matrix

Use Case Comparison

Primary tools per use case. Each cell shows the tool name, its identity in 2–3 lines, and pricing tier.

USE CASE	Anthropic Claude family · claude.ai	Google Gemini family · google.com/gemini
Q&A, writing, general assistant	Sonnet 4.6 default · Opus 4.7 for complex tasks · Haiku 4.5 for speed. Best hallucination reliability at 36% — the industry's lowest at the frontier. Web, iOS, Android. FREE TIER PRO \$20/MO MAX \$100–200/MO	Gemini 3.1 Pro / 3.5 Flash backbone. Deep Search, 1M context window, native Workspace integration (Gmail, Docs, Calendar, Maps). Cheapest paid AI plan at \$7.99/mo for Plus. FREE (FLASH) AI PRO \$19.99/MO
Coding Assistant Code gen, refactoring, debugging, IDE agents	Claude Code Terminal-first IDE agent. 4.2M weekly active developers. SWE-bench Pro 64.3% — highest in the industry. 91% CSAT, NPS 54. Authors 4% of all public GitHub commits (SemiAnalysis, May 2026). PRO \$20 · MAX \$100/\$200 Claude Agent SDK Python + TypeScript SDK for programmable agentic loops. MCP-native (Anthropic invented)	Antigravity 2.0 VS Code fork + Go CLI. Multi-agent orchestration with parallel subagents. Artifacts for verification. Launched May 19, 2026 at Google I/O. Note: Gemini CLI sunset June 18, 2026 for non-enterprise. FREE PREVIEW AI PRO \$19.99/MO Gemini Code Assist Enterprise IDE plugin for VS Code and IntelliJ. Standard and Enterprise

USE CASE	Anthropic Claude family · claude.ai Google Gemini family · google.com/gemini	
	Google ADK and OpenAI Agents SDK. API PER-TOKEN	
Design & UI/UX Prototyping, mockups, design-to-code	Claude Artifacts Interactive React, SVG, HTML and Mermaid prototypes rendered live inside chat. Best-in-class for conversational UI prototyping — build, iterate, export in one session. ALL PAID PLANS Live Artifacts Artifacts that pull live data from MCP connectors — dashboards and trackers that refresh on open, not snapshots. Launched April 2026. Also produces .docx, .pptx, .xlsx natively. ALL PAID PLANS	Stitch Dedicated UI-to-code design tool at stitch.withgoogle.com. Conversational canvas, Voice Canvas mode, Gemini 3-powered. Exports directly to Antigravity via MCP. Upgraded at Google I/O 2026 with Stitch Agent. FREE Imagen 4 Photoreal image generation for mockups, brand assets, UI illustrations. Nano Banana 2 variant provides speed and cost leadership at under 50% of Imagen 4 pricing. API PER-IMAGE
Video Generation Text-to-video, cinematic, ad production	No native capability Anthropic has no video generation offering. Must source video from Google Veo 3.1 or a third-party provider via MCP. N/A	Veo 3.1 Cinematic 1080p/4K with native synced audio, scene extension, and ingredients-to-video pipeline. ~\$0.75/sec via API. Veo 3.1 Lite at under 50% of that. Editorial winner for professional video. ~\$0.75/SEC API AI PRO/ULTRA Flow Veo + Imagen storyboarded production app. Full video editorial workflow, not just clip generation. Included in AI Pro/Ultra subscription. AI PRO \$19.99/MO Gemini Omni Single unified model for text, image, audio, and video input and output. YouTube Shorts integration (10s

USE CASE	Anthropic Claude family · claude.ai Google Gemini family · google.com/gemini	
Image Generation AI art, product photos, marketing visuals	No native capability No raster image generation. Claude can analyze, describe, and iterate on images as input but cannot produce them. N/A	Imagen 4 Highest-fidelity image generation. Photoreal output, fine detail, strong text rendering. Best for marketing, product shoots, and brand assets requiring maximum quality. API PER-IMAGE Nano Banana 2 Speed and cost leader (Gemini 3.1 Flash Image). Under 50% of Imagen 4 pricing. New default for high-volume production image pipelines where speed matters over peak quality. LOW-COST API
Research & Knowledge Document synthesis, web research, RAG	Claude Projects Persistent RAG workspace. MCP connectors to FactSet, S&P Capital IQ, PitchBook, Moody's, Morningstar, LSEG, Chronograph. The only AI stack with governed financial data connectors out of the box. PRO/MAX/TEAM Claude Managed Agents — Outcomes Agents that self-grade against a rubric. Dreaming consolidates memory across sessions. Harvey (legal AI) achieved 6× task-completion lift. Wisedocs cut document review time by 50%. ENTERPRISE BETA	NotebookLM Best-in-class source-grounded RAG. Audio Overviews, Video Overviews, Mind Maps, Slide Decks, Data Tables. 80+ language support. Highest citation fidelity of any research tool. FREE 50 QUERIES/DAY PLUS \$19.99/MO Gemini Deep Research Multi-source structured web research. Integrated in Gemini App as Notebooks (April 2026). Broad web + uploaded document coverage. Good balance of breadth and citation. AI PRO/ULTRA
Browser & Computer Use	Claude Cowork OS-level agent on macOS (post-Vercept acquisition, Feb 2026).	Gemini in Chrome GA for Premium subscribers. Note: Project Mariner was shut down May

USE CASE	Anthropic Claude family · claude.ai Google Gemini family · google.com/gemini	
	native, cross-app workflows. PRO+ Claude for Chrome Browser extension GA for all paid users. DOM-aware page reading, form filling, web task execution. Works in tandem with Cowork for full OS + browser coverage. ALL PAID PLANS	AI PRO/ULTRA
Data Analysis SQL, dashboards, stats, data pipelines	Claude + MCP Connectors MCP servers for BigQuery, Snowflake, PostgreSQL, Redshift. Write and explain SQL, build dashboards, analyze large schemas in a 1M-token context window. Best long-context data reasoning. API PER-TOKEN	BigQuery ML + Gemini in BigQuery Data-warehouse-grounded analytics agents native inside BigQuery. SQL generation, anomaly detection, forecasting — no context window copy-paste needed. Best for GCP-native data stacks. BIGQUERY PRICING NotebookLM — Data Tables Extracts and structures tabular data from documents with source citations. Bridges unstructured document data into structured analysis without writing SQL. FREE TIER
Document Creation Word docs, spreadsheets, presentations	Claude Artifacts (.docx / .pptx / .xlsx) Generates formatted Word, PowerPoint, and Excel files directly in chat. M365 add-ins with MCP-connected financial data (FactSet, S&P, Moody's, LSEG, PitchBook) for finance teams. PAID PLANS	Google Workspace AI Gemini embedded in Docs, Slides, Sheets, and Gmail. AI writing, summarization, Slide generation, formula assistance. NotebookLM can generate structured Slide Decks from source documents. WORKSPACE SUBSCRIPTION

USE CASE	Anthropic Claude family · claude.ai Google Gemini family · google.com/gemini	
On-Device / Mobile AI Offline inference, privacy, zero latency	No on-device strategy Claude iOS and Android apps are thin cloud clients. Anthropic is explicitly cloud-only. No native on-device SDK or offline capability exists. CLOUD-ONLY	Gemini Nano 4 On-device SLM via Android AICore. NPU-accelerated on Pixel, Samsung, Snapdragon, MediaTek. Developer Preview April 2026. Built on Gemma 4. 140+ languages, multimodal input. \$0 inference cost. \$0 INFERENCE AI Edge SDK + ML Kit GenAI APIs Drop-in Android APIs for summarization, smart reply, and rewrite. Production-eligible on AICore-enabled devices. Cross-platform via LiteRT (incl. iOS, Web). FREE SDK
Agentic & Multi-Agent Long-horizon tasks, orchestration, automation pipelines	Claude Managed Agents Hosted agent runtime with sandboxes, vault credentials, MCP Tunnels. Outcomes self-grading against rubric. Dreaming memory consolidates across sessions. Harvey: 6× completion lift. Wisedocs: 50% doc review reduction. PER-TOKEN (PUBLIC BETA) Advisor Tool Pairs a fast executor model with a higher-intelligence advisor model mid-generation — cost-optimized agentic workloads without paying Opus rates for every step. BETA API HEADER	Antigravity 2.0 + ADK 2.0 + Managed Agents API Full stack: Agent Studio visual builder, code-first Python/Go/Java/TS SDK (ADK), Agent Engine managed runtime, CLI. A2A inter-agent protocol. 200+ models in Agent Garden including Claude and GPT-5.5. GCP CONSUMPTION PRICING

USE CASE	Anthropic	Google
	Claude family · claude.ai	Gemini family · google.com/gemini
HIPAA, VPC, SOC 2, data residency, compliance	27 AWS regions. HIPAA-eligible under the existing AWS BAA. VPC endpoints, IAM, CloudTrail, Bedrock Guardrails for PII and policy enforcement. ZDR. Bedrock pricing matches Anthropic direct API rates. HIPAA-ELIGIBLE SOC 2 · ISO 27001	Platform Rebranded from Vertex AI at Cloud Next '26. 200+ models, data residency in 10 countries, HIPAA BAA via GCP, CMEK. Agent Engine at \$0.0864/vCPU-hr. Broadest multi-region data residency footprint. HIPAA-ELIGIBLE 10-COUNTRY RESIDENCY
	Claude Enterprise SAML/OIDC SSO, SCIM, audit logs, ZDR, HIPAA BAA, BYOK launching H1 2026. Self-serve Enterprise introduced in 2026. Private network routing via Bedrock or Vertex. CUSTOM PRICING	

Performance vs Efficiency — The Car Analogy

AI models trade benchmark power against token cost and hallucination rate — exactly like cars trading horsepower against fuel economy.

THE FERRARI Peak performance, expensive per mile. You buy it for the lap time, not the tank range. Use for the hardest jobs where accuracy is non-negotiable. Expect to pay.	THE PORSCHE / BMW High performance, still daily-driveable. Excellent power-to-cost ratio. The right choice for 80% of real production workloads.
THE TOYOTA / HONDA Optimised for mileage. Lower peak output but high reliability per dollar. Route your high-volume, low-complexity traffic here for 30–60% cost reduction.	THE LAMBORGHINI Raw horsepower in some dimensions, but burns through fuel and misfires often. Brilliant on the track it was built for. Terrible on roads that require consistency.

Claude Opus 4.7 Peak performance · premium cost · industry-best accuracy	Claude Sonnet 4.6 High performance · daily driveable · production default
---	--

Reliability (inverse hallucination)	64% reliable — 36% halluc.	Reliability (inverse hallucination)	Excellent — near Opus level
Cost efficiency (lower = cheaper)	\$5 in / \$25 out per Mtok	Cost efficiency	\$3 in / \$15 out per Mtok
SWE-bench Pro (coding accuracy)	64.3% — #1 in industry	SWE-bench Verified	79.6%
Context window	1M tokens	Speed vs Opus	~2× faster
Fuel warning: New tokenizer inflates effective cost up to 35% vs prior Opus. Gate behind cost guardrails; only route here when accuracy is the primary constraint.		Recommended default for 80% of production workloads. Best balance of accuracy, speed, and cost in the Anthropic family.	

Claude Haiku 4.5	Toyota Prius	Gemini 3.1 Pro Preview	BMW M3
Maximum mileage · routing & classification specialist		Best price-to-reasoning ratio · still Preview status	
Reliability	Excellent	Reliability (inverse hallucination)	50% reliable — 50% halluc.
Cost efficiency	\$1 in / \$5 out per Mtok	Cost efficiency	\$2 in / \$12 out per Mtok
Best use case	Routing, classification, extraction	GPQA Diamond (pure reasoning)	94.3% — #1 in industry
Speed	Fastest in Anthropic family	Context window	1M (2M practical)
Bulk-route here: Intelligent routing to Haiku for low-complexity tasks cuts inference cost 30–60% in mixed workloads (CloudZero, 2026 research).		Half the cost of Opus 4.7 with better pure reasoning. Caveat: still "Preview" label as of May 2026 — don't anchor production SLAs to it yet.	

Gemini 3.5 Flash	Honda Civic Type R	GPT-5.5	Lamborghini Huracán
Sporty performance · economy running cost · new default		Dominant on specific tracks · high cost · high fuel-misfire rate	
Reliability	Moderate (similar to 3.1 Pro)	Reliability (inverse hallucination)	14% reliable — 86% halluc.
Cost efficiency	\$1.50 in / \$9 out per Mtok	Cost efficiency	\$5 in / \$30 out per Mtok
MCP Atlas Score	83.6% — beats 3.1 Pro on coding	Terminal-Bench 2.0	82.7% — #1 for terminal agents
Speed vs 3.1 Pro	~4× faster	OSWorld (computer use)	78.7% — #1 for computer use
New price-performance leader for high-volume production. Launched May 19, 2026 — replaces Flash as the Google default recommendation.		86% hallucination risk: Use only where grounding guardrails exist. Wins terminal and computer-use benchmarks by a wide margin — deploy there and nowhere else without RAG.	

GPT-5.5 Pro

Formula 1 Car

Extreme performance · extreme cost · one specific race

Reliability **Better than GPT-5.5 standard**

Cost efficiency **\$30 in / \$180 out per Mtok**

BrowseComp (web research) **90.1% — best web research**

Best for **Frontier research, hard math**

6× more expensive than GPT-5.5 standard. Reserved for BrowseComp-level research or AIME-grade math where nothing else suffices.

Methodology: Hallucination rates from Artificial Analysis AA-Omniscience benchmark — which rewards models that decline to answer when uncertain. GPT-5.5's 86% reflects answering when wrong rather than declining. Real-world impact depends on whether your application layers grounding and citation guardrails. Treat single benchmarks as directional, not decisive. Pricing as of May 2026 — OpenAI doubled GPT-5.5 vs GPT-5.4 at launch; Google removed Pro-tier free API access April 1, 2026. Re-validate budget assumptions quarterly.

Token Cost & Licence Reference

Per-million-token pricing as of May 2026. Lower input cost = better mileage. All values USD.

MODEL	VENDOR	INPUT \$/MTOK	OUTPUT \$/MTOK	CACHED \$/MTOK	CONTEXT	HIPAA-ELIGIBLE
Claude Sonnet 4.6	Anthropic	\$3.00	\$15.00	\$0.30	1M	Yes (Bedrock)
Claude Haiku 4.5	Anthropic	\$1.00	\$5.00	—	200K	Yes (Bedrock)
Gemini 3.1 Pro Preview	Google	\$2.00	\$12.00	\$0.20	1M (2M prac.)	Yes (Vertex)
Gemini 3.5 Flash	Google	\$1.50	\$9.00	\$0.15	1M	Yes (Vertex)

MODEL	VENDOR	INPUT \$/MTOK	OUTPUT \$/MTOK	CACHED \$/MTOK	CONTEXT	HIPAA- ELIGIBLE
GPT-5.5 Pro	OpenAI	\$30.00	\$180.00	—	1M	Enterprise only

Pricing notes: Gemini 3.1 Pro charges \$4/\$18 above 200K tokens. GPT-5.5 adds a 10% regional premium for non-US endpoints and 2× input / 1.5× output above 272K tokens. Claude Opus 4.7 new tokenizer can inflate effective cost up to 35% vs prior Opus. Cached token pricing applies to prompt caching features. HIPAA eligibility applies to specific deployment paths (Bedrock BAA, Vertex GCP BAA) — not to consumer tiers or preview products.

Dominant Vendor by Layer

Where each vendor leads as of May 2026, and who trails closest.

LAYER	DOMINANT	RUNNER-UP
Pure reasoning at frontier price	GOOGLE Gemini 3.1 Pro — GPQA 94.3% at \$2/Mtok	Anthropic (Opus 4.7)
Complex multi-file coding accuracy	ANTHROPIC Opus 4.7 — SWE-bench Pro 64.3%	OpenAI (Codex)
Terminal & agentic computer use	OPENAI GPT-5.5 + Codex — Terminal-Bench 82.7%, OSWorld 78.7%	Anthropic (Cowork)
Browser automation (shipped today)	ANTHROPIC Claude Cowork — 72.5% OSWorld	OpenAI (Atlas)
On-device mobile inference	GOOGLE Gemini Nano 4 + AI Edge SDK — only production on-device stack	None (OpenAI + Anthropic: cloud-only)

Video generation	GOOGLE Veo 3.1 / Gemini Omni	OpenAI (Sora 2 sunseting Sep 2026)
Image generation	GOOGLE Imagen 4 / Nano Banana 2	OpenAI (ChatGPT Images 2.0)
Design → code	GOOGLE Stitch + Antigravity 2.0	Anthropic (Artifacts)
Research — document grounded	GOOGLE NotebookLM — best citation fidelity	Anthropic (Projects + MCP)
Research — web breadth	OPENAI ChatGPT Deep Research — broadest coverage	Google (Gemini Deep Research)
Hallucination reliability	ANTHROPIC Opus 4.7 — 36% AA-Omniscience (best)	Google (50%); OpenAI distant (86%)
Token cost efficiency (frontier)	GOOGLE Gemini 3.5 Flash — \$1.50/\$9 per Mtok	Anthropic (Haiku 4.5 at \$1/\$5 for lighter tasks)
Long context (1M+)	GOOGLE Gemini 3.1 Pro — 2M practical	Anthropic (Sonnet/Opus 4.x — 1M GA)
MCP ecosystem maturity	ANTHROPIC Inventor of MCP — deepest tooling, agent SDK	Google (Antigravity ADK), OpenAI (Secure MCP Tunnel)
Enterprise compliance breadth	TIED Anthropic on Bedrock + Google Vertex (both HIPAA, ZDR, ISO)	OpenAI (Azure, 10 regions, EKM)
Prompt injection robustness	ANTHROPIC Opus 4.5: 4.7% attack success	Google (~12.5%); OpenAI (~21.9%)

(industry
best)

Source: All data from Gaurav Arora's AI-Resources-Comparison repo (May 21, 2026), cross-referenced against Artificial Analysis, SemiAnalysis, Anthropic/Google/OpenAI official benchmarks, Presenc AI developer survey (April 2026), and JetBrains AI Pulse Survey (January 2026, n=10,000+ professional developers).